

Hyperparameter Optimization

Milana M. Wolff

May 07, 2024

1 Introduction

In this assignment, we optimize relevant hyperparameters for a small selection of classification models using the wine quality dataset. This widely used dataset contains a variety of physicochemical input features, such as wine density and acidity, along with expert ratings for red Vinho Verde wines. We approach the hyperparameter optimization problem by selecting three classifiers with competitive performance under default hyperparameter configurations: Ridge, bagging, and random forest. We then conduct hyperparameter optimization with a Bayesian optimization approach and nested resampling.

2 Dataset Description

The dataset used for this assignment contains physicochemical quantitative input features and sensory quantitative output features (i.e., an expert wine score) for the red variant of the Portuguese "Vinho Verde" wine [1]. The dataset includes 1599 observations and eleven input features, including fixed acidity, volatile acidity, citric acid, residual sugar, chlorides, free sulfur dioxide, total sulfur dioxide, density, pH, sulphates, and alcohol content. According to the UC Irvine Machine Learning Repository website, "the classes are ordered and not balanced (e.g. there are many more normal wines than excellent or poor ones)", with a total of 1319 observations rated as 5 or 6 and a mere 28 observations rated with the highest and lowest scores (3 and 8) [2]. This robust dataset includes no missing values to be imputed. We use the eleven listed input features to predict the wine quality measurement as the target during hyperparameter optimization.

3 Experimental Setup

We use a Python implementation running on the `teto-gpu` partition of the Beartooth cluster (default Python version 3.7.16 (GCC 11.2.0)), with an allocation of four nodes, 16 cores per node, and 64 GB of memory, as two of

the classifiers do not complete hyperparameter optimization when executed on Colaboratory, where we initially attempted to run this code. After importing the `scikit-learn` package, we use Ridge, bagging, and random forest classifier models.

We optimize the alpha, tolerance, solver, and maximum iteration hyperparameters for the Ridge classifier models. These hyperparameters control regularization strength, the precision of the solution, the solver used in computational routines to obtain the best solution, and the maximum number of iterations for conjugate gradient solver, respectively [3]. For the bagging classifier, we use the maximum samples (as a proportion of total samples), number of estimators, and maximum feature hyperparameters [4]. We use the number of estimators, splitting criterion, and maximum tree depth for the random forest classifier models [5].

The ranges and priors for numerical hyperparameters and the options for categorical hyperparameters are tabulated below.

Classifier	Ridge	Bagging	Random Forest
Numerical Hyperparameter Ranges	alpha: (1.0, 5.0) tol: (0.0001, 1.0) max_iter: (100, 100000)	n_estimators: (100, 10000) max_samples: (0.01, 1.0) max_features: (0.01, 1.0)	n_estimators: (100, 10000) max_depth: (2, 10)
Numerical Hyperparameter Priors	log-uniform (all)	log-uniform (all)	log-uniform (all)
Categorical Hyperparameters	'svd', 'cholesky', 'lsqr', 'sparse_cg'		'gini', 'entropy', 'log_loss'

We load the wine data using the `pandas` library and use all eleven features for classification without additional pre-processing steps.

To perform hyperparameter optimization, we use Bayesian cross-validation, in which we specify a search space and use the default 3 folds. We implement this approach with the approach with the `BayesSearchCV()` method provided by `scikit-learn`. These 3 folds constitute the inner folds for nested resampling. We perform 50 iterations, sampling 10 points in parallel and executing 10 jobs. The relatively low number of hyperparameter combinations evaluated is due to time constraints on performing a search on the hyperparameter space. We use balanced accuracy as a scoring criterion in the Bayesian search to account for the underlying class imbalance problem in the wine quality dataset. The Bayesian CV is performed within a call to `cross-validate` with the default 5 folds, which constitute the outer folds for nested resampling.

4 Results

We report the average balanced accuracy scores across the 5 folds of the outer cross-validation and the standard deviation for these scores. The hyperparameters yielding the best performance across the outer CV folds were selected as the optimal hyperparameters. For comparison, we also report the average balanced accuracy score for 5-fold CV using the default hyperparameters for each model type. These results are tabulated for the Ridge, bagging, and random forest classifiers below.

Classifier	Ridge	Bagging	Random Forest
Generalization Score (Optimized)	0.225 \pm 0.024	0.278 \pm 0.038	0.270 \pm 0.032
Best (Optimized)	0.2464	0.3363	0.3070
Generalization Score (Default Hyperparameter Settings)	0.247 \pm 0.027	0.340 \pm 0.042	0.351 \pm 0.043

The hyperparameter configuration for the best estimator for the Ridge classifier in the search space defined was: `alpha=3.9832`, `max_iter=87432`, `solver='svd'`, and `tol=0.0001`. This configuration achieved the highest performance in the outer CV, with a balanced accuracy score of 0.2464.

In comparison, the default hyperparameters for the Ridge classifier are `alpha=1.0`, the default maximum number of iterations is determined by `scipy.sparse.linalg`, the solver is automatically chosen to fit the data, and `tol=0.0001`.

For the bagging classifier, the best hyperparameter configuration obtained used `max_samples=0.3456`, `n_estimators=100`, and the default value for `max_features` (the maximum number of features used for classification, expressed as a fraction of the total number of features). This configuration achieved a balanced accuracy score of 0.3363 on the outer CV.

In comparison, the default hyperparameter settings are `max_features = 1.0` (i.e., all features), `max_samples = 1.0` (all samples), and `n_estimators = 10`. The default estimator is a decision tree, as the bagging classifier is an ensemble model.

The random forest classifier with the best performance in nested resampling used `criterion='entropy'`, `max_depth=6`, `n_estimators=104` as the hyperparameter configuration. A balanced accuracy score of 0.3070 was achieved with these settings.

The default hyperparameter settings are `max_depth = None` (i.e., no limit on maximum tree depth), `criterion = 'gini'`, and `n_estimators = 100`.

Compared to default hyperparameter settings, the optimized hyperparameter settings did not improve performance, and yielded slightly worse results in all cases. This may be the result of a relatively low number of iterations of the Bayesian search CV (due to memory and time constraints) and achieving local

minima or, in the specific case of the random forest classifier, limiting the maximum tree depth of the search space too greatly to achieve optimal performance. In all other cases, the default hyperparameter settings were configurations possible within the Bayesian search space.

5 Code

<https://github.com/COSC5557/hyperparameter-optimization-mwolff2021-1>

References

- [1] In: (). URL: <http://www.vinhoverde.pt/en/>.
- [2] Paulo Cortez, A. Cerdeira, F. Almeida, et al. “Wine Quality”. In: (2009). DOI: <https://doi.org/10.24432/C56S3T>.
- [3] URL: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.RidgeClassifier.html#sklearn.linear_model.RidgeClassifier.
- [4] URL: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.BaggingClassifier.html#sklearn.ensemble.BaggingClassifier>.
- [5] URL: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>.